

POC1 — Azure Shopping List

Finance Procure-to-Pay POC. Azure-only consumption. Third-party services (HeyGen, external SaaS) and the Custom Control Plane UI are costed separately.

Shopping list

#	Resource	Tier / config	Purpose
1	Azure AI Foundry — Agent Service	Hosted Agents (preview), 2 deployments × 2 replicas	Hosts Finance Agent + Fleet Manager. No runtime charge — pay only for models and tools.
2	Azure OpenAI — GPT-5.4-mini	PAYG	Primary model for all agent executors and Fleet Manager batch reasoning.
3	Azure OpenAI — GPT-5.4 (full, non-Pro)	PAYG	~10% escalation for complex exception classification + hourly Fleet Manager deep summaries.
4	Azure Functions — Elastic Premium EP1	1 vCPU / 3.5 GB, always-on	Hosts Durable Functions orchestrations (workflow envelope, HITL waits, checkpoint/replay).
5	Azure API Management — Standard v2	1 unit	AI Gateway: token rate limits, content safety, semantic cache, load balancing, MCP governance.
6	Azure API Center	Preview (free)	Registry for models, MCP tools, skills.
7	Azure Cosmos DB — Serverless (NoSQL)	Serverless account, 2-3 containers	Workflow state + append-only action ledger.
8	Azure Blob Storage — Hot LRS	Standard GPv2, ~200 GB	DF Task Hub + audit archive.
9	Azure Key Vault	Standard	MCP credentials (Workday, D365 F&O, Maconomy).
10	Azure AI Search — Basic	1 SU	Substrate for Foundry IQ (vendor master, purchasing policy, tax-rule corpora).
11	Azure Document Intelligence	PAYG, prebuilt invoice model	Deterministic OCR; LLM only on low-confidence fields.
12	Azure Container Apps — Consumption	0.25 vCPU / 0.5 GiB × 3 apps	Hosts MCP servers (Workday, D365 F&O, Maconomy).
13	Azure Monitor — Log Analytics + App Insights	PAYG Analytics tier	OTEL from Durable Functions, MAF, GHCP SDK sessions, APIM, MCP servers.
14	Azure Event Grid	Standard topic	Telemetry fan-out to Fleet Manager.

15	Azure Content Safety / Foundry Guardrails	PAYG per text record	Four-point interception on agent I/O (input, output, tool-call, tool-response).
16	Microsoft Entra Agent ID	Existing Entra P2 tenancy	3 agent identities; no incremental cost.

Assumptions

Assumption	Value	Rationale
Fixed infrastructure operating mode	24 × 7	APIM, Durable Functions host, AI Search, Container Apps cannot be paused meaningfully in a development environment.
Fleet Manager operating mode	24 × 7 from day 1	Reasons over telemetry in ~3-minute batches with an hourly deep pass on GPT-5.4.
Invoices processed end-to-end	10,000	Dev iterations + test + demo + slack.
Pages per invoice (OCR)	2	One invoice + one PO attachment on average.
Agent executor hit rate	~30% of invoices	P2P is deterministic-by-default; LLM only fires on exception paths (low-confidence OCR, GL coding, reconciliation exceptions).
Test iteration multiplier	3×	Same invoice reprocessed across dev iterations, debugging, demo rehearsals.
Agent executor token profile	8k input / 1k output per invocation	Invoice content + skill prompt + MCP tool responses.
Fleet Manager batch rate	~500 passes/day (every ~3 minutes) + hourly deep pass	Aggregated reasoning over telemetry, not per-event.
Fleet Manager context profile	5k in / 500 out per batch; 8k in / 2k out for hourly deep	Telemetry summaries, not raw event dumps.
Content Safety inspection	4 passes × ~20k agent invocations	Matches Foundry Guardrails' four intervention points.

Out of scope (deliberately excluded)

Excluded	Reason
Voice / ACS / GPT-Realtime / avatars	POC2 scope.

Custom Control Plane UI + SignalR + React hosting	Costed separately.
Microsoft Fabric capacity + Fabric IQ	Not load-bearing for POC1; cost-centre / agency hierarchy can be sourced directly from D365 F&O or Dataverse. Re-appears in POC2 where episodic memory and cross-entity ontology are genuinely needed.
Agent 365 / Purview DLP / Defender for Cloud AI licences	Absorbed by existing WPP M365 E5 + Defender tenancy.
Private Endpoints / VNet / NAT gateway	Not baselined; adds ~\$135 per 8 weeks if compliance posture mandates.
Secondary region / DR	POC runs in a single region.
Third-party SaaS (Workday, D365 sandbox, Maconomy licences)	WPP-provided per POC spec.
Production capacity (500 concurrent × 30 markets)	Different order of magnitude — separate exercise.

Model family

Frontier-only. No legacy 4-series models.

- **GPT-5.4-mini** — default for agent executors and Fleet Manager batch reasoning.
- **GPT-5.4** (full, non-Pro) — escalation path for complex exceptions and hourly Fleet Manager deep summaries.
- **GPT-5.4 Pro** — not used in POC1; reserved for scenarios requiring maximum reasoning depth.