

POC2 — Azure Shopping List

HR Talent Lifecycle POC. 12-week "frontier" sprint. Azure-only consumption. Third-party services (HeyGen avatar, Greenhouse / LinkedIn / Workday SaaS licences) and the Custom Control Plane UI are costed separately.

Shopping list

#	Resource	Tier / config	Purpose
1	Azure AI Foundry — Agent Service	Hosted Agents (preview), 3 deployments × 2-3 replicas	Hosts Hiring Agent + Fleet Manager + External A2A Candidate Agent bridge. No runtime charge — pay only for models and tools.
2	Azure OpenAI — GPT-5.4-mini	PAYG	Primary model for agent executors (job design, CV triage, interview coordination, onboarding) and Fleet Manager batch reasoning.
3	Azure OpenAI — GPT-5.4 (full, non-Pro)	PAYG	Compliance narrative (jurisdiction-aware right-to-work, GDPR, EU AI Act interpretation), offer personalisation, Fleet Manager hourly deep pass, complex exception handling.
4	Azure OpenAI — GPT-Realtime	PAYG	Speech-to-speech front end for voice screening.
5	Azure Communication Services	PAYG voice calling + PSTN	Telephony substrate for candidate voice screening calls.
6	Azure Functions — Elastic Premium EP1	1 vCPU / 3.5 GB, always-on	Hosts Durable Functions orchestrations for the long-running (12-week) hiring workflow — phase boundaries, HITL waits at zero compute, timer escalation, checkpoint/replay.
7	Azure API Management — Standard v2	1 unit	AI Gateway: token rate limits, content safety, semantic cache, model load balancing, MCP governance, A2A governance (preview) for external candidate agent.
8	Azure API Center	Preview (free)	Registry for models, MCP tools, skills, A2A agent cards. Jurisdiction-specific skill promotion (USA/DE) through lifecycle gates.
9	Azure Cosmos DB — Serverless (NoSQL)	Serverless account, multiple containers	Workflow state, action ledger, and episodic memory store (levelling history, past-hire outcomes for Fleet Manager precedent retrieval).
10	Azure Blob Storage — Hot LRS	Standard GPv2	Candidate artefacts (CVs, transcripts, screening recordings), offer letters, DF Task Hub, audit archive.
11	Azure Key Vault	Standard	MCP credentials (Greenhouse, LinkedIn, Workday, MS Graph, ServiceNow, ACS).

12	Azure AI Search — Standard S1	1 replica / 1 partition	Substrate for Foundry IQ across jurisdiction-specific corpora: USA employment law, German BetrVG, GDPR consent guidance, EU AI Act classifier corpus, WPP people handbooks, comp benchmarking data. Standard tier needed for corpus size + permission filtering.
13	Microsoft Fabric — F2 capacity	PAYG, shared attribution	Substrate for Fabric IQ: headcount ontology, agency hierarchy, comp bands, cross-entity matrix navigation (role in Media Agency, budget from WPP Corp), levelling history. Load-bearing for episodic memory and budget reasoning.
14	Azure Container Apps — Consumption	0.25 vCPU / 0.5 GiB × 6 apps	Hosts MCP servers: Greenhouse ATS, LinkedIn Recruiter, Workday (hiring), MS Graph, ServiceNow, ACS control plane.
15	Azure Monitor — Log Analytics + App Insights	PAYG Analytics tier	OTEL across 10+ agent roles, Durable Functions orchestration, MAF workflow executors, APIM, MCP servers. Per-phase cost attribution.
16	Azure Event Grid	Standard topic	Telemetry fan-out to Fleet Manager from multi-agent, multi-phase hiring workflow. Higher event volume than POC1.
17	Azure Content Safety / Foundry Guardrails	PAYG per text record	Four-point interception on agent I/O. Task Adherence detects drift from jurisdiction policy in the system message.
18	Azure AI Foundry — Evaluators	PAYG	500-CV synthetic evaluation run (task adherence, tool call accuracy, bias markers, groundedness). Continuous evaluation on production traces.
19	Azure Document Intelligence	PAYG, prebuilt / layout model	Light use — right-to-work document parsing, offer letter extraction. Much lower volume than POC1.
20	Microsoft Entra Agent ID	Existing Entra P2 tenancy	3 agent identities (Hiring Agent, Fleet Manager, External A2A bridge); no incremental cost.
21	Work IQ MCP	Existing M365 tenancy (preview)	M365 work-graph grounding for Personal Agents, Interview Coordinator (timezone / calendar), org-topology escalation routing. No separate Azure line.

Assumptions

Assumption	Value	Rationale
Build operating mode	Dev-style utilisation curve — not 24×7 production	Weeks 1-4 mostly local dev; weeks 5-8 integration; weeks 9-12 demo prep + multi-jurisdiction rehearsals.

Fleet Manager operating mode	Activated from week 4 onwards, 24×7 during weeks 4-12	Realistic once workflows start existing in volume.
End-to-end hiring workflows executed	~50 (test + demo + rehearsal)	Conservative; 15-20 concurrent at demo peak × multiple rehearsals.
Synthetic CV evaluation	500 CVs × 3-5 eval runs	POC2 spec requires 500-CV eval via Foundry Evaluators.
Voice screening calls	~50 calls × 15 minutes	Screening agent demo runs, multi-jurisdiction variants.
Jurisdictions exercised	USA + Germany	POC2 frontier requirement — jurisdiction switching via APIM routing + jurisdiction-specific skills + Foundry IQ corpora.
Agent roles loaded as skills	10+ skills on the Hiring Agent Hosted Agent	Per solution \$4.4 — skills specialise within one domain-scoped agent, not separate agents.
Agent executor token profile	8k-15k input / 1k-2k output per invocation	Larger than POC1 — CVs are longer than invoices; compliance narratives are token-heavy.
Fleet Manager profile	Batch reasoning every ~3 min + hourly deep pass on GPT-5.4	Higher event rate than POC1 due to 10+ agent roles per workflow.
Episodic memory reads	Fabric IQ + Work IQ on every major phase transition	Load-bearing for "recall past hires levelled too low" demo.
A2A external candidate agent	1 integration, APIM-governed (preview)	Single demo flow to show external-agent interop.

Out of scope (deliberately excluded)

Excluded	Reason
HeyGen avatar generation	Third-party; user excluded external parties.
Candidate web portal hosting	Platform / Control Plane shopping list.
Custom Control Plane UI + SignalR + React hosting	Platform / Control Plane shopping list.
Agent 365 / Purview DLP / Defender for Cloud AI licences	Absorbed by existing WPP M365 E5 + Defender tenancy.
M365 Copilot + M365 Agents SDK surface licences	Existing WPP M365 tenancy.
Third-party SaaS (Greenhouse, LinkedIn Recruiter, Workday sandbox licences)	WPP-provided per POC spec.
Private Endpoints / VNet / NAT gateway	Add if compliance posture mandates network isolation.

Secondary region / DR	POC runs in a single region (EU for DE jurisdiction, US for USA).
Production run-rate (500 concurrent × 30 markets)	Different order of magnitude — separate exercise.
MAI-Voice-1 (preview TTS)	Primary voice uses GPT-Realtime (GA); MAI-Voice-1 additive only.

Model family

Frontier-only. No legacy 4-series models.

- **GPT-5.4-mini** — default for agent executors (job design, CV triage, interview coordination, onboarding) and Fleet Manager batch reasoning.
- **GPT-5.4** (full, non-Pro) — compliance narrative (jurisdiction-aware right-to-work, GDPR, EU AI Act), offer personalisation, Fleet Manager hourly deep pass, complex exception handling.
- **GPT-Realtime** — speech-to-speech for voice screening.
- **GPT-5.4 Pro** — not used in POC2; reserved for scenarios requiring maximum reasoning depth.

What differs from POC1

- **Voice stack** added: GPT-Realtime + ACS for candidate screening.
- **Fabric capacity** is load-bearing (episodic memory + cross-entity ontology), unlike POC1 where it was descoped.
- **AI Search Standard S1** instead of Basic — jurisdiction corpora (USA, DE, GDPR, EU AI Act) exceed Basic tier capacity.
- **A2A governance** in APIM (preview) — for external candidate agent demo.
- **Foundry Evaluators** — 500-CV synthetic eval requires this line.
- **More Container Apps** (6 MCP servers vs 3 in POC1).
- **Document Intelligence** minimal / optional — hiring has light document OCR needs vs P2P invoice-heavy.
- **Work IQ MCP** is load-bearing for Personal Agents and Interview Coordinator (calendar, timezone, org topology).