Technical Response Sections [Scott + Artur]

These sections replace §4–§13 and §16–§17 in the vendor response document. They are grounded in the validated solution architecture. Every product claim is referenced.

---

4. Framework Architecture

4.1 Microsoft Agent Stack Overview

Our architecture expresses the **determinism ↔ agentic spectrum** explicitly across three cooperating execution layers, with a Fleet Manager observing from above, a central governance layer, and a custom Control Plane UI.

Deterministic by default, agentic by exception. WPP's processes contain both rule-driven steps (three-way match, jurisdiction routing, payment file generation) and reasoning-driven steps (CV triage, voice screening, compliance narrative review). Mixing them cleanly requires a layered substrate:

Fleet Managers: always-on GHCP SDK Hosted Agents on Azure AI Foundry. Domain-scoped (hiring, finance, compliance). Consume telemetry from all workflows via Azure Event Grid. Reason about fleet health, SLA risk, anomalies. Compose the exception queue for the Control Plane. Push assessments via SignalR.

Layer 1 — Durable envelope: Azure Durable Functions (GA). One orchestration instance per workflow. Owns phase boundaries, HITL waits at zero compute (days or weeks), timer escalation, parallel fan-out/fan-in, checkpoint/replay with geo-replicated state in Azure Storage. Fully deterministic — code-defined, event-sourced replay.

Layer 2 — Workflow graph: Microsoft Agent Framework (MAF) workflows (v1.0 GA) connected to DF via the durable task extension — Microsoft's Durable Agent Orchestration pattern (Feb 2026). Each workflow phase is a graph of typed executors: plain-function executors for deterministic operations, agent executors where LLM reasoning is required, and validator executors between them (judge/executor separation expressed as two nodes and an edge). Pregel BSP execution. MAF's own pause/resume and checkpointing are preserved across DF replay.

Layer 3 — Agent executor contents: ephemeral GHCP SDK sessions on Foundry Hosted Agents. Invoked only from MAF agent executor nodes. Each session loads skills and MCP tools for the task, reasons about it, calls tools through hooks, writes state to Cosmos DB, emits OTEL telemetry, and returns a typed result to its MAF executor.

Skill crystallisation is the evolution path: proven agent executors graduate from LLM-generated to deterministic code, versioned as skills in API Center, and swapped into the MAF graph as plain-function executors — agent executor preserved as fallback for exceptions.

Central Governance sits alongside all three tiers:

- **Azure API Center + APIM AI Gateway** (GA): one cohesive control point for everything addressable by an agent — models, MCP tools, A2A agents, skills, and APIs. **API Center** is the unified registry: lifecycle management (Design → Preview → Production → Deprecated), GitHub Actions sync, cross-cloud discovery (Azure / GCP / AWS / on-prem). **APIM AI Gateway** is the runtime: model load balancing, failover, spillover, token rate limits, per-team / per-workflow budget control, semantic caching, jurisdiction-based routing; MCP tool governance (auth injection, rate limiting, content safety); A2A governance (preview); REST→MCP auto-generation. One pane, one policy engine, one audit trail.
- **Foundry Control Plane** (GA): agent fleet inventory, model registry (1900+ models), guardrails (PII, Task Adherence), built-in evaluators (quality, safety, drift).
- **Agent 365 + Entra** (GA May 2026): Entra Agent ID (first-class agent identity, not service accounts), agent registry, lifecycle management (activate, block, delete), Conditional Access, Purview DLP, Defender.

Intelligence Layer — agents need three kinds of grounding (enterprise documents, business semantics, work context). Microsoft has a purpose-built product for each. All three are MCP-addressable and sit behind APIM AI Gateway alongside any other tool, governed identically:

- **Foundry IQ** (preview): unified agentic retrieval over enterprise corpora. Self-reflective query planner with configurable retrieval reasoning effort. Federates SharePoint, Fabric, OneLake, Blob, AI Search, web, and MCP behind one permission-aware endpoint. Built on Azure AI Search. Used by Compliance Agent (employment law, GDPR, EU AI Act), Job Design (comp benchmarking), Skill Amplification (Fleet Manager surfacing precedents to operator).
- **Fabric IQ** (preview): semantic intelligence layer over WPP's Fabric / OneLake estate. Business ontology, semantic model, graph engine for multi-hop reasoning. Used by Budget Agent (headcount, cost-centre, agency hierarchy), ROI reporting, cross-entity matrix navigation.
- **Work IQ MCP** (preview): M365 work graph + memory layer (collaboration patterns, calendar, expertise). Used by Personal Agents, Interview Coordinator (timezone, availability), Org Topology / Escalation routing.

4.2 Agent Framework vs Agent Surfaces

Layer	What It Is	Components
Framework	The agent runtime, workflow graph engine, durable orchestration envelope, state store, governance, tool integration, knowledge grounding, and observability.	GHCP SDK (agent runtime), Microsoft Agent Framework (workflow graph), Foundry Hosted Agents, Durable Functions (durable envelope), APIM AI Gateway + API Center, Foundry IQ + Fabric IQ + Work IQ (Intelligence Layer), Foundry Control Plane, Agent 365 + Entra, Cosmos DB, Application Insights, Log Analytics
Surfaces	The channels through which humans interact with agents and agent outputs.	M365 Copilot (Teams), Email (Adaptive Cards), Custom Control Plane UI (React), Web Portal, Voice (ACS + GPT-Realtime), ServiceNow

A single domain agent (e.g. Hiring Agent) is surface-agnostic. The same agent executors inside its MAF workflows produce outputs that the human's Personal Agent (PA) delivers to whatever surface that human uses. The PA is a GHCP SDK agent surfaced in M365 Copilot via M365 Agents SDK. Every human has one. Humans make all decisions. The PA prepares and executes.

4.3 Control Plane Architecture

Two layers.

Foundry Control Plane (existing product, not custom build): fleet health dashboards, agent inventory, model registry, guardrails configuration, continuous evaluation, Defender and Purview integration. This covers WPP Refs 8.1, 8.5, 8.14, 22.3.

Custom Control Plane UI (React application powered by Fleet Manager agents): this is where WPP's operator-specific requirements are met. No vendor ships what WPP is asking for — a fleet management interface where 1 human governs 20–50 concurrent workflows with intelligent exception surfacing. This is custom build.

Capability	What it does	WPP Ref
Fleet Dashboard	Fleet Manager composes workflow-level status, SLA tracking, agency/market/jurisdiction filtering. Not raw telemetry — the Fleet Manager's assessment.	31.1

Exception-Only Queue	Of N active workflows, the operator sees only the 2–3% needing attention. Fleet Manager composes this based on business impact, confidence, SLA urgency.	31.2
Instant Situational Awareness	Click any workflow → what happened, why it stopped, what was tried, what the Fleet Manager recommends, options. Pre-composed by Fleet Manager. <5 seconds.	31.3
Bulk HITL	Batch similar decisions (8 interview schedules, all low-risk). Approve in one action. Raises events on all waiting Durable Functions instances simultaneously.	31.4
Autonomy Dials	Per-workflow threshold sliders. Writes to config store, effective immediately, audit-logged.	21.1
Skill Amplification	Fleet Manager proactively surfaces policy, precedents, recommended approach when operator is uncertain.	31.5
Role-Based Views	HR BP sees hiring. Finance BP sees budget gates. Entra RBAC.	10.1
Cost Dashboard	Per-workflow cost attribution from OTEL + APIM token metrics.	26.4

Data sources for the custom UI: Application Insights APIs (traces, cost), Foundry REST APIs (agent inventory), APIM metrics (token consumption), workflow state store in Cosmos DB (phase status, approvals, action ledger), Fleet Manager assessments (SignalR real-time).

4.4 Multi-Agent Orchestration

We do not use separate agent processes for each specialist role. Instead, skills provide specialisation within agent executors in a MAF workflow graph, per domain.

A domain-scoped Hosted Agent (e.g. Hiring Agent, `hiring-agent@wpp`) runs ephemeral GHCP SDK sessions from MAF agent executors, each loading different skills depending on the workflow phase. Each skill defines its own role, allowed tools, model assignment, and governance rules. The outcome — specialised capabilities with distinct roles and model assignments — matches WPP's requirement for heterogeneous teams. The architectural advantages over separate agents: no inter-agent communication overhead, shared workflow context without message passing, simpler governance (one Entra identity per domain), easier operationalisation at fleet scale.

Two coordination substrates, layered:

- **Within a phase — MAF workflow graph**: Pregel BSP execution, typed edges, fan-out/fan-in, conditional routing, validator executors, pause/resume. Stable MAF orchestration patterns (sequential, concurrent, handoff, group chat, Magentic-One) cover the multi-agent topologies WPP requires.
- **Across phases — Azure Durable Functions**: long-running envelope, HITL waits at zero compute (days/weeks), timer escalation, checkpoint/replay, geo-replicated state. Invokes each phase's MAF workflow as a durable activity via the MAF durable task extension.

Both substrates adapt based on runtime data — not static DAGs. Supported topologies end-to-end: sequential (interview after screening), parallel fan-out/fan-in (sourcing + job design concurrently), conditional (compliance flag triggers additional review branch in the MAF graph), timer escalation, bulk HITL.

---

5. Core Capability Mapping

5.1 Core Framework Capabilities

Capability	Solution	Status
Multi-agent orchestration	Two-layer: Durable Functions (across phases) + MAF workflows (within phase) + skills-based specialisation in GHCP SDK agent executors	GA (DF, MAF v1.0, MAF Durable Task extension), Tech Preview (GHCP SDK)
Stateful workflow management	Cosmos DB (workflow state, action ledger) + Durable Functions (orchestration state in Azure Storage) + MAF workflow checkpointing preserved across DF replay	GA
Tool/function integration	MCP servers governed through APIM. REST-to-MCP gateway auto-generates tool definitions from OpenAPI specs.	GA (APIM), MCP preview
Connectors	MCP servers for Workday, Greenhouse, LinkedIn, D365 F&O, Maconomy, ServiceNow, MS Graph, ACS, HeyGen — all via APIM	Custom build per connector
Human-in-the-loop	GHCP SDK hooks intercept non-revocable actions inside agent executors. MAF validator executors enforce structural/policy checks before non-revocable executors fire. Durable Functions `wait_for_external_event` at zero compute for long waits. PA surfaces decisions to humans. Agent 365 routes to right person.	GA (Durable Functions, MAF), custom (HITL flow)
Durable execution	Durable Functions checkpoint/replay envelope + MAF native pause/resume preserved across DF replay (via MAF Durable Task extension). Cosmos DB continuous backup with PITR. Geo-replicated state.	GA
Memory	Cosmos DB (facts + episodic), skills (procedural), Foundry IQ (semantic retrieval), Fabric IQ (business ontology), Work IQ (M365 episodic), GHCP SDK session (working)	GA (data stores), Preview (IQ products), custom (memory patterns)
Model-agnostic	GHCP SDK works with any model from Foundry catalog (1900+). Per-skill model assignment. APIM routes and governs.	GA
Control Plane	Foundry Control Plane (platform) + Custom Control Plane UI (operator experience)	GA (Foundry) + custom build

5.2 Desirable Capabilities

Capability	Solution	Status
Rollback / compensating actions	Action ledger tracks revocable vs non-revocable. Hooks block non-revocable before execution. Durable Functions triggers compensating actions on rollback.	Custom build (proven pattern)
Self-healing	Durable Functions retry with backoff. MAF workflow conditional routing to recovery branches. APIM circuit breakers and model failover. Fleet Manager routes exceptions to humans.	GA

Multi-surface engagement	PA delivers to Teams, Email, Control Plane UI, ServiceNow, Web, Voice — all from the same agent executor output	GA (M365 Agents SDK) + custom (Control Plane UI, voice)
A2A with off-platform agents	APIM A2A governance: AgentCards, JSON-RPC, SSE	Preview
Evaluation framework	Foundry Evaluators: task adherence, tool call accuracy, safety, groundedness, sensitive data exposure. Continuous evaluation on production traffic.	GA
Workflow crystallisation	Agent executors in the MAF graph graduate from generative (LLM-driven) to deterministic (plain-function) as patterns mature. Skills versioned through API Center lifecycle gates (Design → Production). Agent executor preserved as exception fallback.	Custom build

5.3 Advanced Capabilities

Capability	Solution	Status
Voice screening	GPT-Realtime (speech-to-speech) + ACS Call Automation. MAI-Voice-1 (TTS, preview).	GA (GPT-Realtime, ACS), Preview (MAI-Voice-1)
Avatar onboarding video	Agent executor drafts script; downstream deterministic executor calls HeyGen API MCP to produce branded video.	Custom MCP server
Fine-tuning	Azure AI Foundry fine-tuning. We prioritise skill crystallisation over model fine-tuning.	GA
Knowledge extraction	Threadlight accelerator: creates skills from interviews and unstructured data.	Built, demonstrated

---

6. Governance, Security and Compliance

6.1 Agent Identity and Governance

Each Hosted Agent has a dedicated Entra Agent ID — a first-class identity in Entra ID, distinct from service accounts. Domain-scoped:

Hosted Agent	Entra Agent ID	Tool Access
Hiring Agent	`hiring-agent@wpp`	Greenhouse, LinkedIn, Workday (hiring), Graph, ACS
Finance Agent	`finance-agent@wpp`	Workday (finance), D365 F&O, Maconomy
Fleet Manager	`fleet-manager@wpp`	Read-only: telemetry, state store. No downstream tool invocation.

Identity lives on the container, not on individual sessions. When a human triggers an action, the loop acts on-behalf-of that human. For autonomous phases, the loop uses the Hosted Agent's app-only identity. Every action is attributed to the correct identity in the audit trail.

Agent 365 provides lifecycle management (activate, block, delete), Conditional Access for agents, and integration with Purview (DLP, audit) and Defender (threat detection).

6.2 Observability and OpenTelemetry

Three layers of observability:

1. **Foundry Tracing**: every GHCP SDK session (inside a MAF agent executor) emits OTEL spans via GHCP SDK OTEL TracerProvider; MAF workflow execution emits executor-lifecycle spans natively; DF emits orchestration spans. Model calls, tool calls, tokens, latency, cost. Spans carry workflow ID, phase, jurisdiction, model, token count.
2. **APIM metrics**: token usage, latency, errors per model/tool/agent. Central cost tracking.
3. **Fleet Manager assessment**: event-driven reasoning over telemetry. Fleet health, anomaly detection, SLA risk, exception prioritisation. This is the default Control Plane view.

Cost attribution is per workflow, per phase, per model, per consumer. Visible in both Foundry dashboards and the custom Control Plane UI.

6.3 Data Protection and Compliance

Three enforcement layers. The LLM cannot bypass any of them.

Foundry Guardrails (preview): intercepts at four intervention points — input, output, tool-call, tool-response. PII detection. Task Adherence detects policy drift.

APIM AI Gateway (GA): jurisdiction-based model routing (DE workflow → EU endpoint only). Content safety policy. Token rate limiting. All MCP tool calls governed.

Agent 365 + Entra (GA May 2026): per-agent RBAC on downstream resources. Conditional Access. Purview DLP on agent interactions. Defender for threat detection.

Policy-as-code: APIM policies via APIOps CI/CD. Foundry Guardrails per agent. Jurisdiction-specific skills in SKILL.md files. All Git-committed, PR-reviewable, audit-trailed.

Jurisdiction switching: workflow state carries `jurisdiction`. APIM routes to region-appropriate model endpoint. Jurisdiction-specific skills (right-to-work, works council, GDPR consent) load automatically based on workflow context. Adding a new jurisdiction = adding new skill files, not modifying agent code.

Audit: Azure Log Analytics with 7–12 year retention. Immutability enforced via Azure Storage export with immutability policies. Every tool call, model call, enforcement decision, and human interaction logged. Queryable via KQL and Microsoft Sentinel.

---

7. Protocol Support

Protocol	Status	Implementation
MCP	GA	Primary integration pattern. GHCP SDK natively supports MCP. All enterprise systems exposed as MCP servers. APIM provides REST-to-MCP gateway (auto-generates tool definitions from OpenAPI specs). MCP governance via APIM.
A2A	Preview	APIM A2A governance: AgentCards, JSON-RPC task lifecycle, SSE streaming. Agents are both A2A clients and servers. A2A is not required for core architecture — used in POC2 for external candidate agent demo.
OpenTelemetry	GA	Native throughout. GHCP SDK OTEL + Foundry Tracing.

8. Enterprise Integration

8.1 LOB Application Integration

All enterprise systems are integrated via MCP servers governed through APIM AI Gateway. Auth complexity is abstracted at the platform level — agent developers never manage tokens. Azure Key Vault stores credentials. APIM injects auth at request time. Automated rotation without agent downtime.

System	Auth Method	MCP Integration
Workday	SAML-bridged via Okta	MCP server, APIM-governed
LinkedIn Recruiter	OAuth 2.0 AuthCode	MCP server, APIM-governed
Greenhouse ATS	API App Credentials	MCP server, APIM-governed
Microsoft Graph	OBO flow	MCP server, APIM-governed
ServiceNow	API key	MCP server, APIM-governed
Dynamics 365 F&O	Native	MCP server, APIM-governed
Maconomy	REST adapter (APIM REST-to-MCP gateway)	Auto-generated from OpenAPI spec
ACS (Voice)	Managed identity	MCP server
HeyGen	API key	MCP server
Citizen-dev integrations (SharePoint, Outlook, Dataverse, etc.)	Logic App connectors	Azure Logic Apps workflow exposed as MCP tool via APIM REST→MCP gateway. 1,400+ prebuilt connectors. No-code path for new tools.

Supported OAuth 2.0 grant types: Authorization Code, Client Credentials, SAML-bridged (Okta), PKCE, Device Flow, OBO. All via Entra External ID federation with Okta as primary IdP.

9. Development Experience

Mode	Persona	Solution
Pro-code	Platform engineers, full-stack developers	GHCP SDK Python (primary) + TypeScript. Full access to skills, MCP, hooks, OTEL instrumentation, model selection. MIT open-source.
Low-code agents	Citizen developers	Copilot Studio visual designer for conversational agents and simpler workflows. Supported via Agent 365 (not Foundry Hosted Agents). Custom Control Plane UI natively supports Copilot Studio agents. Not recommended for complex autonomous workflows.

Low-code MCP tools	IT teams adding tools without writing Python	Azure Logic Apps visual workflows chaining 1,400+ prebuilt connectors. Logic App exposed as MCP tool via APIM REST→MCP gateway. Governed identically to hand-written MCP servers.
Low-code config	Operators, process owners	Custom Control Plane UI: skill library (browse, fork, customise), tool catalog, governance rules, autonomy dials, template management. No code required.
Agentic builder	Domain experts	A MAF agent executor generates SKILL.md files from natural language specifications. Registered in API Center. Human reviews and approves. Built and demonstrated.
Knowledge extraction	Transitioning staff, SMEs	Threadlight accelerator: creates skills from interviews and unstructured data. Produces machine-actionable SKILL.md files.

All agent artefacts are Git-committable: skills as SKILL.md, MCP server code, APIM policies as code, MAF workflow definitions (Python/.NET), Durable Functions orchestrations, autonomy thresholds. Low-code artefacts (Copilot Studio) register through Agent 365 alongside pro-code artefacts.

---

10. Code-as-Truth: Auditability and Transparency

Every artefact is version-controlled and inspectable:

- **Skills**: SKILL.md files in Git. Registered in Azure API Center with lifecycle management (Design → Preview → Production → Deprecated). GitHub Actions integration for syncing from repositories.
- **APIM policies**: version-controlled via APIOps CI/CD. PR review gates. Per-environment policy sets.
- **Orchestrations**: Durable Functions code in Git. CI/CD deployment with environment-specific configuration.
- **Governance rules**: autonomy thresholds, compliance rules, jurisdiction skills — all in version-controlled configuration stores with change tracking (who, when, why).
- **Traceability**: every routing decision, model selection, and tool call is fully traceable in OTEL spans and audit logs. Every entry in the action ledger links to the OTEL span that produced it, including the reasoning chain and tool call that triggered it.

A team of 5 developers manages 50 agents across dev, staging, and production using CI/CD pipelines with PR review gates. Non-technical auditors inspect agent authorisations, actions, and rationale via the Foundry Control Plane compliance dashboard and Log Analytics KQL queries.

---

11. Non-Functional Requirements Response

NFR	Target	How
Availability	99.9% per region	Azure SLA-backed across all services
RTO	<5 minutes	Azure Traffic Manager + Cosmos DB automatic failover + Durable Functions replay
RPO	Near-zero	Cosmos DB continuous backup with PITR. Durable Functions checkpoints in geo-redundant Azure Storage.
Control Plane latency	<5 seconds	Fleet Manager pre-composes situational context. SignalR push.

Concurrent workflows (pilot)	5,000+	Durable Functions auto-scaling. Cosmos DB horizontal partitioning. Multiple Hosted Agent deployments.
Concurrent workflows (prod)	50,000+	Same architecture, scaled. Tiered model usage + crystallisation reduce inference bottleneck. Reference architecture to be provided separately.
Audit log retention	7–12 years	Azure Log Analytics archive tier. Immutability via Azure Storage export.
Data residency	Platform-level	APIM routes by jurisdiction. Cosmos DB regional deployment. Log Analytics regional workspaces. Not developer-configured.
Workflow state	Survives full restart	Durable Functions replays from checkpoint. Cosmos DB serves from geo-replica. Maximum loss: current in-flight phase replays from start.

---

12. POC 1: Finance Intelligent Procure-to-Pay

Duration: 8-week sprint. **Operator**: Finance Controller via Control Plane. **Concurrency**: 30–50 concurrent invoice workflows.

Hosted Agent: Finance Agent (`finance-agent@wpp`).

Skills: Intake/OCR (Azure Document Intelligence), Validation (three-way match, duplicate detection), Routing (GL coding, cost centre allocation), Approval (threshold-based routing with HITL gates), Payment (payment file generation), Reconciliation (statement matching, exception identification).

MCP integrations: Workday, Dynamics 365 F&O, Maconomy — all governed through APIM AI Gateway. Sandbox/mock APIs provided by WPP.

Execution shape: a Durable Functions orchestration per invoice. Each phase (intake, validation, routing, approval, payment, reconciliation) is a MAF workflow graph. Most executors are plain functions; agent executors are used only where genuine reasoning is needed; validator executors sit between agent output and any downstream non-revocable action.

Demonstrates:

- Deterministic-by-default MAF workflow graph (three-way match, payment file generation, routing are plain functions — no LLM)
- Agent executors limited to low-confidence OCR extraction, GL coding reasoning, exception classification — each followed by a validator executor
- Multi-phase orchestration via Durable Functions, with each phase invoked as a durable activity wrapping a MAF workflow
- HITL approval gates (Finance BP interacts via Adaptive Card in Outlook, routed through PA; DF `wait_for_external_event` at zero compute)
- Bulk approval for batched low-risk items
- Rollback and compensating transactions (non-revocable actions gated by GHCP SDK hooks inside agent executors AND by MAF validator executors)
- Fleet Manager monitoring 30–50 concurrent workflows
- Exception-only Control Plane view
- OTEL cost attribution per invoice across all three layers
- Foundry Guardrails inside agent executors (PII detection, content safety)
- Full audit trail from receipt to payment

- Mid-workflow platform restart with DF replay and MAF workflow checkpoint resume

---

13. POC 2: People Talent Lifecycle Agent Team

Duration: 12-week sprint following POC 1. **Operator**: HR Business Partner (London) via Control Plane. **Concurrency**: 15–20 concurrent hiring workflows. **Scenario**: Hire a Senior Data Engineer at a WPP agency in USA. **Five human participants**: Hiring Manager (LA, Teams), HR BP (London, Control Plane), Finance BP (Mumbai, email), IT Ops (Chennai, ServiceNow), Candidate (external, web + voice).

Hosted Agent: Hiring Agent (`hiring-agent@wpp`).

Skills: Budget & Approvals, Job Design, Sourcing, Triage, Screening, Interview Coordinator, Compliance (jurisdiction-aware), Offer, Onboarding, Voice Screening.

MCP integrations: Greenhouse ATS, LinkedIn Recruiter, Workday (hiring), Microsoft Graph (calendar/email), ServiceNow (IT provisioning), Azure Communication Services (voice), HeyGen (avatar) — all governed through APIM AI Gateway.

Execution shape: a Durable Functions orchestration per hire. The ~10 phases (budget & approvals, job design, sourcing, CV triage, voice screening, interview coordination, compliance, offer, JML onboarding, avatar welcome) are each MAF workflow graphs. Deterministic executors dominate (sourcing queries, interview scheduling, JML provisioning tickets, offer letter templating, HeyGen video generation). Agent executors are bounded — JD drafting, CV scoring, voice screening, compliance narrative, offer personalisation. Validator executors enforce structural and policy checks between agent output and downstream action.

Demonstrates (all POC 1 capabilities plus):

- Layered orchestration: DF envelope across 12 weeks + MAF workflow graphs per phase + GHCP SDK sessions in agent executors
- Deterministic-by-default graphs annotated per phase (see solution.md §15) — transparently shows what is code vs what is LLM
- Voice screening with structured scoring (GPT-Realtime + ACS) followed by validator executor
- CV parsing with crystallisation pipeline (agent CV scorer → deterministic classifier in API Center, agent preserved as fallback)
- Episodic memory (recall past hires levelled too low) via workflow state store + Fabric IQ
- A2A interop with external candidate agent (APIM-governed)
- Jurisdiction-aware compliance (USA vs Germany enforcement switching via APIM routing + jurisdiction-specific skills + Foundry Guardrails + MAF validator executors for GDPR consent, EU AI Act classification)
- Autonomy dials (configurable auto-shortlist thresholds, adjustable at runtime — controls whether a phase routes to agent executor or falls back to deterministic baseline)
- Skill amplification (Fleet Manager surfaces policy + precedents)
- Process evolution (Fleet Manager proposes crystallisation candidates after completed workflows)
- Synthetic CV evaluation (500 CVs via Foundry Evaluators for bias/accuracy testing)
- Avatar onboarding video (HeyGen API MCP)
- Threadlight knowledge extraction demo (interview HR SME, produce executable skills in the MAF graph)
- 5 humans across 4 timezones interacting simultaneously via different surfaces

---

16. Customer References

Microsoft will provide with the full submission:

1. Written testimonials from 3 enterprise customers using Microsoft's agentic stack at comparable scale and complexity to WPP.
2. Reference contacts: 3 customers available for direct calls, each with Engineering Lead, AI/ML Ops Lead, and Head of Transformation (SVP+).

Reference customers to be confirmed upon mutual NDA execution. Sectors: media, financial services, professional services.

---

17. Portability and Exit Strategy

Lock-in mitigation at every layer:

Layer	Portability
Agent logic	GHCP SDK is open-source (MIT). Skills are SKILL.md markdown files. WPP retains ability to self-host, fork, or migrate.
Models	GHCP SDK works with any model from the Foundry catalog (1900+ including OpenAI, Anthropic, Google, Meta, Mistral, open-source). Switching models = configuration change in APIM.
Tool layer	MCP is an open standard. All enterprise integrations are MCP servers, portable to any MCP-compliant platform.
Interop	A2A is an open standard. Agent definitions using A2A are portable.
Telemetry	OpenTelemetry is an open standard. Traces portable to any OTEL-compatible backend.
State	Workflow state is JSON in Cosmos DB, exportable via standard APIs. Durable Functions state in Azure Storage, extractable.
Orchestration	Durable Functions is Azure-specific. The pattern (event-driven workflow orchestration) is reproducible on Temporal, AWS Step Functions, etc.
Policies	APIM policies exportable. Governance rules are in version-controlled configuration.

Exit strategy: export skills (Git), export MCP servers (code), export state (Cosmos DB export), export policies (APIM export). The agentic logic is in portable artefacts.

Appendix: Known Constraints

Constraint	Impact	Mitigation
GHCP SDK in tech preview	API surface may change	Core patterns (skills, MCP, hooks) proven in production. MIT open-source.
Microsoft Agent Framework v1.0 (Oct 2025)	Framework is young	Core runtime and workflows are GA. MAF durable task extension for Azure Functions is productised by Microsoft as the Durable Agent Orchestration pattern. Orchestration patterns stable. Fallback: GHCP SDK + DF combination works without MAF — MAF adds the deterministic graph primitive.

Foundry Hosted Agents: max 5 replicas (preview)	Scaling ceiling	Multiple deployments, or Azure Container Apps with Foundry telemetry.
Foundry Guardrails tool-call interception (preview)	May not be GA for POC	GHCP SDK session hooks provide equivalent enforcement. Guardrails are additive.
APIM A2A governance (preview)	A2A features maturing	Not required for core architecture. HTTP gateway primitives work today.
API Center skill registry (preview)	No native Git sync	Uses GitHub Actions workflows. Core skill execution is GHCP SDK native.
GHCP SDK + Foundry Hosted Agents integration	Hosting adapter needs custom work	Primary integration engineering task for the POC.
Agent 365 GA: May 2026	Integration with Hosted Agents unclear	Entra Agent ID usable independently. Needs POC validation.
Foundry IQ / Fabric IQ / Work IQ in public preview	Intelligence Layer products are new; APIs evolving	All three are MCP-addressable — fall back to direct Azure AI Search + Fabric SQL + Graph API if needed. The IQ products are an upgrade path, not a single point of failure.
MAI-Voice-1 (preview)	No SLA	GPT-Realtime (GA) is the primary voice path. MAI-Voice-1 is additive.
Copilot Studio on Foundry Hosted Agents	Not supported	Copilot Studio agents supported via Agent 365. Control Plane UI supports both.

References

#	Source
1	GHCP SDK repository
2	GHCP SDK skills
3	GHCP SDK MCP support
4	GHCP SDK hooks
5	GHCP SDK OTEL
6	GHCP SDK agent loop
7	Azure AI Foundry Hosted Agents

8	Foundry Control Plane fleet monitoring
9	Foundry Guardrails
10	Foundry Evaluators
11	Foundry model catalog
12	APIM AI Gateway capabilities
13	APIM MCP server support
14	APIM REST-to-MCP gateway
15	APIM A2A governance
16	APIOps CI/CD
17	Agent 365 overview
18	Entra Agent ID
19	Entra Agent OBO flow
20	Conditional Access for agents
21	Purview for Agent 365
22	Defender AI threat protection
23	Azure Durable Functions
24	Cosmos DB continuous backup
25	Azure Log Analytics retention
26	Azure API Center
27	M365 Agents SDK
28	GPT-Realtime audio
29	ACS Call Automation
30	Azure Document Intelligence
31	Foundry IQ overview
32	Foundry IQ — connect agents
33	Fabric IQ overview
34	Work IQ MCP overview
35	Azure Logic Apps
36	Microsoft Agent Framework overview
37	MAF Workflows

38	MAF Workflow Executors
39	MAF Durable Task Extension for Azure Functions
40	MAF Durable Agent Orchestration tutorial
41	MAF v1.0 release announcement
42	Building Human-in-the-Loop AI Workflows with MAF